

Sequence Based T Cell Epitope Prediction: Verification and Improvements

Alexander Crowell, Jack Holland, and Vineetha Paruchuri

Abstract—Successful prediction of T cell epitopes is very useful in vaccine design and immunotherapy. Obviously, computational prediction is a lot cheaper than experimental validation for T cell epitope prediction. Researchers have previously predicted the T cell epitopes with varying degrees of success. In this paper, we use Support Vector Machine, a machine learning algorithm, to do the same, and we compare our results to pan-specific approaches of T cell epitope prediction, like NetMHCII, NetMHCIIpan, and the MultiRTA methods. Our analysis showed that ;insert data;

Keywords—T Cell Epitopes, Epitope Mapping, Epitope Prediction, NetMHC, MultiRTA, Neural Networks, Support Vector Machines, Machine Learning.

I. INTRODUCTION

TCELL epitopes are the part of an antigen, usually a specific peptide sequence, which is bound by the Major Histocompatibility Complex (MHC). Two primary groupings of MHC proteins, MHC I and MHC II, are respectively present either on all cells or only antigen presenting cells. Due to differences in the structure of their binding pockets, T cell epitopes presented by MHC class I molecules are peptides between 8 and 11 amino acids in length, whereas those presented by MHC class II are 13-17 amino acids in length.

The high degree of polymorphism exhibited by MHC II alleles means that experimental validation of epitope regions is cost intensive. Hence, the computational methods for accurate prediction of binding peptides to class II molecules is a topic of interest. Researchers have previously tried to computationally predict T cell epitopes, with various degrees of success, using Artificial Neural Networks, such as in NetMHCII [?] and NetMHCIIpan [?], or Position Specific Scoring Matric [?] such as in TEPITOPEpan [?], or Position Weight Matrix, such as in PREDIVAC [?].

In this project, we use the machine learning technique of Support Vector Machines to predict T cell epitopes, and compare the performance of our approach to allele-specific and pan-specific approaches like NetMHC, NetMHCIIpan, and MultiRTA [?]. The results show that ;insert something about our performance;

II. NETMHCII

NetMHCII is a neural network based approach for MHC class II predictions. It uses a novel artificial neural network-based method, *NN-align*, which allows for simultaneous identification of the MHC class II binding core and binding affinity. As mentioned in [?], the NN-align method includes explicit encoding of the peptide flanking residues in terms of amino acid composition and length, as well as a novel scheme for

neural network training that deals with the data redundancy inherent in the peptide data due to multiple examples of identical binding cores. The NN-align method is trained on a large data set of more than 14,000 quantitative peptide MHC binding values covering 14 HLA-DR alleles.

A. NetMHCII Implementation

NN-align was implemented as a conventional feed-forward artificial neural network method. The method consists of a two-step procedure that simultaneously estimates the optimal peptide binding register (core) and network weight configuration. The score of a 9 mer peptide was calculated using the conventional feed-forward algorithm. The network weights were updated using gradient descent back-propagation. The predictive performance of the method was significantly improved by incorporation of information from residues flanking the peptide-binding core.

B. NetMHCII Performance, and Metrics Used

All statistical comparisons were made using one-tailed binomial tests. For each comparison, it is calculated how often one method outperforms the other (excluding ties), and based on these numbers, one-tailed p-values were calculated. P-values less than 0.05 were taken to be significant. Performance across different methods is calculated using AUC.

C. Inferences from NetMHCII

Our critique goes here. I'll update once we discuss what we wanna say about it. Maybe something like it is comprehensive, and tested for all these things x y z, which is why we think it is a good method for MHCII binding predictions despite ;cite limitations;, so we want to compare ours to this etc.?

III. SUPPORT VECTOR MACHINES

Support Vector Machine is a supervised learning algorithm used for classification and regression analysis - they identify patterns in the data. Once we train the SVM on some data (called the training data set), it classifies new examples into one of two classes, thereby making it a non-probabilistic binary linear classifier. An SVM can also perform non-linear classification using *kernel trick*, which just means that it doesn't need to compute the coordinates of data in a high-dimensional space, but rather just the inner products between images of all pairs of data, thereby making it computationally cheaper.

Given the fact that it is a binary linear classifier that also performs non-linear classification, we felt that, as a first step, using SVM as the choice of machine learning algorithm made sense.

A. Formal Definition

Given some training data D , a set of n points of the form

$$D = \{(x_i, y_i) \mid x_i \in R^p, y_i \in \{-1, 1\}\}_{i=1}^n$$

where the y_i is either 1 or -1, indicating the class to which the point x_i belongs. Each x_i is a p -dimensional real vector. We want to find the maximum-margin hyperplane that divides the points having $y_i = 1$ from those having $y_i = -1$. Any hyperplane can be written as the set of points x satisfying

$$\mathbf{w} \cdot x - b = 0,$$

where \cdot denotes the dot product and w the (not necessarily normalized) normal vector to the hyperplane. The parameter $\frac{b}{\|\mathbf{w}\|}$ determines the offset of the hyperplane from the origin along the normal vector \mathbf{w} .

If the training data are linearly separable, we can select two hyperplanes in a way that they separate the data and there are no points between them, and then try to maximize their distance. The region bounded by them is called "the margin". These hyperplanes can be described by the equations

$$\mathbf{w} \cdot \mathbf{x} - b = 1$$

and

$$\mathbf{w} \cdot \mathbf{x} - b = -1$$

By using geometry, we find the distance between these two hyperplanes is $\frac{2}{\|\mathbf{w}\|}$, so we want to minimize $\|\mathbf{w}\|$. As we also have to prevent data points from falling into the margin, we add the following constraint:

for each i either

$$\mathbf{w} \cdot \mathbf{x}_i - b \geq 1 \quad \text{for } \mathbf{x}_i \text{ of the first class or } \mathbf{w} \cdot \mathbf{x}_i - b \leq -1 \quad \text{for } \mathbf{x}_i \text{ of the second.}$$

This can be rewritten as:

$$y_i(\mathbf{w} \cdot \mathbf{x}_i - b) \geq 1, \quad \text{for all } 1 \leq i \leq n. \quad (1)$$

We can put this together to get the optimization problem:

$$\text{Minimize } (\|\mathbf{w}\|, b) \quad \text{subject to } (for any } i = 1, \dots, n) y_i(\mathbf{w} \cdot \mathbf{x}_i - b) \geq 1.$$

IV. METHODS

This section provides a brief overview of how we implemented the algorithm and how we validated it.

A. Datasets

We took the peptides from IEDB, and subdivided them according to allele type. We also tried out other data repositories, like the Dana-Farber [?] repository, but since all entries in it have a positive binding value, we didn't find it suitable for our problem. We also ran our algorithm on the MultiRTA [?] dataset. Data selection and sanitization is a crucial step here, because any irregularities in data can affect the learning of the SVM in unpredictable ways.

B. Implementation of the SVM and Metrics

We implemented SVM in python using scikit-learn [?], and used scikit-learn functions such as roc_auc_score etc. for metrics.

C. Encoding Peptides

D. AUC as a Comparison Metric

Area Under the Curve is used as a metric throughout, because other pan-specific approaches used comparable metrics. (Maybe include some more detail?)

blah blah blah

V. RESULTS

A. Individual Alleles

B. Combined Alleles

C. Results Compared to RTA

D. Results Compared to NetMHCII

E. Results Compared to NetMHCIIpan

Can describe what results we got in cross validation, what we compared to, how it scaled up etc.

VI. CONCLUSION

Our inferences in how this performed and what it means etc.

APPENDIX A DIVISION OF WORK

A. Alexander Crowell

B. Jack Holland

C. Vineetha Paruchuri

REFERENCES

- [1] H. Kopka and P. W. Daly, *A Guide to L^AT_EX*, 3rd ed. Harlow, England: Addison-Wesley, 1999.
- [2] Sturniolo T, Bono E, Ding J, Raddrizzani L, Tuereci O, et al. *Generation of tissue-specific and promiscuous HLA ligand database using DNA microarrays and virtual HLA class II matrices*, 7th ed. Nat Biotechnol, 1999, pp.555561